

Charu C. Aggarwal
IBM T J Watson Research Center
Yorktown Heights, NY

Basic Principles of Reinforcement Learning [Motivating Deep Learning]

Neural Networks and Deep Learning, Springer, 2018
Chapter 8.1–8.3

The Complexity of Human Intelligence is Quite Simple!

- Herbert Simon's ant hypothesis:

“Human beings, viewed as behaving systems, are quite simple. The apparent complexity of our behavior over time is largely a reflection of the complexity of the environment in which we find ourselves.”

- Humans are simple, because they are reward-driven entities.
- All of biological intelligence is owed to this simple fact.
- Reinforcement learning attempts to simplify the learning of complex behaviors by using *reward-driven trial and error*.

When to Use Reinforcement Learning?

- Systems that are simple to judge but hard to specify.
- Easy to use trial-and-error to generate data.
 - Video games (e.g., Atari)
 - Board and card games (e.g., chess, Go, Poker)
 - Robot locomotion and visuomotor skills
 - Self-driving cars
- *Reinforcement learning is the gateway to general forms of artificial intelligence!*

Why Don't We have General Forms of Artificial Intelligence Yet?

- Reinforcement learning requires large amounts of data (generated by trial and error).
 - Possible to generate lots of data in some game-centric settings, but not other real-life settings.
- Biological reinforcement learning settings include some unsupervised learning.
 - The number of synapses in our brain is larger than the number of seconds we live!
 - There must be some unsupervised learning going on continuously ⇒ We haven't mastered that art yet.
- Recent results do show promise for the future.

Simplest Reinforcement Learning Setting: Multi-armed Bandits

- Imagine a gambler in a casino faced with 2 slot machines.
- Each trial costs the gambler \$1, but pays \$100 with some unknown (low) probability.
- The gambler suspects that one slot machine is better than the other.
- What would be the optimal strategy to play the slot machines, assuming that the gambler's suspicion is correct?
- **Stateless Model:** Environment at every time-stamp is identical (although knowledge of *agent* improves).

Observations

- Playing both slot machines alternately helps the gambler learn about their payoff (over time).
 - However, it is wasteful *exploration*!
 - Gambler wants to *exploit* winner as soon as possible.
- Trade-off between exploration and exploitation \Rightarrow Hallmark of reinforcement learning

Naïve Algorithm

- **Exploration:** Play each slot machine for a fixed number of trials.
- **Exploitation:** Play the winner forever.
 - Might require a large number of trials to robustly estimate the winner.
 - If we use too few trials, we might actually play the poorer slot machine forever.

ϵ -Greedy Strategy

- Probabilistically merge exploration and exploitation.
- Play a random machine with probability ϵ , and play the machine with highest current payoff with probability $1 - \epsilon$.
- Main challenge in picking the proper value of $\epsilon \Rightarrow$ Decides trade-off between exploration and exploitation.
- **Annealing:** Start with large values of ϵ and reduce slowly.

Upper Bounding: The Optimistic Gambler!

- Upper-bounding represents optimism towards unseen machines ⇒ Encourages exploration.
- Empirically estimate mean μ_i and standard deviation σ_i of payoff of the i th machine using its n_i trials.
- Pick the slot machine with largest value of mean plus confidence interval $= \mu_i + K \cdot \sigma_i / \sqrt{n_i}$
 - Note the $\sqrt{n_i}$ in the denominator, because it is *sample* standard deviation.
 - Rarely played slot machines more likely to be picked because of optimism.
 - Value of K decides trade-off between exploration and exploitation.


Multi-Armed Bandits versus Classical Reinforcement Learning

- Multi-armed bandits is the simplest form of reinforcement learning.
- The model is stateless, because the environment is identical at each time-stamp.
- Same action is optimal for each time-stamp.
 - Not true for classical reinforcement learning like Go, chess, robot locomotion, or video games.
 - *State* of the environment matters!

Markov Decision Process (MDP): Examples from Four Settings

- *Agent*: Mouse, chess player, gambler, robot
- *Environment*: maze, chess rules, slot machines, virtual test bed for robot
- *State*: Position in maze, chess board position, unchanged, robot joints
- *Actions*: Turn in maze, move in chess, pulling a slot machine, robot making step
- *Rewards*: cheese for mouse, winning chess game, payoff of slot machine, virtual robot reward

The Basic Framework of Reinforcement Learning



1. AGENT (MOUSE) TAKES AN ACTION a_t (LEFT TURN IN MAZE) FROM STATE (POSITION) s_t
2. ENVIRONMENT GIVES MOUSE REWARD r_t (CHEESE/NO CHEESE)
3. THE STATE OF AGENT IS CHANGED TO s_{t+1}
4. MOUSE'S NEURONS UPDATE SYNAPTIC WEIGHTS BASED ON WHETHER ACTION EARNED CHEESE



OVERALL: AGENT LEARNS OVER TIME TO TAKE STATE-SENSITIVE ACTIONS THAT EARN REWARDS

- The biological and AI frameworks are similar.
- MDP represented as $s_0 a_0 r_0 s_1 a_1 r_1 \dots s_n a_n r_n$

Examples of Markov Decision Process

- *Game of tic-tac-toe, chess, or Go:* The state is the position of the board at any point, and the actions correspond to the moves made by the agent. The reward is +1, 0, or -1 (depending on win, draw, or loss), which is received at the end of the game.
- *Robot locomotion:* The state corresponds to the current configuration of robot joints and its position. The actions correspond to the torques applied to robot joints. The reward at each time stamp is a function of whether the robot stays upright and the amount of forward movement.
- *Self-driving car:* The states correspond to the sensor inputs from the car, and the actions correspond to the steering, acceleration, and braking choices. The reward is a function of car progress and safety.

Role of Traditional Reinforcement Learning



- **Traditional reinforcement learning:** Learn through trial-and-error the *long-term* value of each state.
- Long-term values are not the same as rewards.
 - Rewards not realized immediately because of stochasticity (e.g., slot machine) or delay (board game).


Reinforcement Learning for Tic-Tac-Toe

- Main difference from multi-armed bandits is that we need to learn the long-term rewards for each action in each *state*.
- An eventual victory earns a reward from $\{+1, 0, -1\}$ with delay.
- A move occurring r moves earlier than the game termination earns *discounted* rewards of $\{\gamma^{r-1}, 0, -\gamma^{r-1}\}$.
 - Future rewards would be less certain in a replay.
- Assume that a fixed pool of humans is available as opponents to train the system (self-play possible).

Generalizing ϵ -Greedy to Tic-Tac-Toe

- Maintain table of values of state-action pairs (initialize to small random values).
 - In multi-armed bandits, we only had values on actions.
 - Table contains unnormalized total reward for each state-action pair \Rightarrow Normalize to average reward.
- Use ϵ -greedy algorithm with *normalized* table values to simulate moves and create a game.
- **After game:** Increment at most 9 entries in the unnormalized table with values from $\{\gamma^{r-1}, 0, -\gamma^{r-1}\}$ for r moves to termination and win/loss.
- Repeat the steps above.

At the End of Training



- Typical examples of normalized values of moves
- ϵ -greedy will learn the strategic values of moves.
- Rather than state-action-value triplets, we can equivalently learn state-value pairs.

Where Does Deep Learning Fit In?

- The tic-tac-toe approach is a glorified “learning by rote” algorithm.
- Works only for toy settings with few states.
 - Number of board positions in chess is huge.
 - Need to be able to *generalize* to unseen states.
 - **Function Approximator:** Rather than a table of state-value pairs, we can have a *neural network* that maps states to values.
 - The *parameters* in the neural network substitute for the table.


Strawman ϵ -Greedy Algorithm with Deep Learning for Chess [Primitive: Don't Try It!]

- Convolutional neural network takes board position as input and produces position value as output.
- Use ϵ -greedy algorithm on output values to simulate a full game.
- **After game of X moves:** Create X training points with board position as input feature map and targets from $\{\gamma^{r-1}, 0, -\gamma^{r-1}\}$ depending on move number and win/loss.
- Update neural network with these X training points.
- Repeat the steps above.

Reinforcement Learning in Chess and Go

- The reinforcement learning systems, *AlphaGo* and *AlphaZero*, have been designed for chess, Go, and shogi.
- Combines various advanced deep learning methods and Monte Carlo tree search.
- Plays positionally and sometimes makes sacrifices (much like a human).
 - Neural network encodes evaluation function learned from trial and error.
 - More complex and subtle than hand-crafted evaluation functions by conventional chess software.

Examples of Two Positions from Alpha Zero Games vs Stockfish



- Generalize to unseen states in training.
- Deep learner can recognize subtle positional factors because of trial-and-error experience with feature engineering.

Other Challenges

- Chess and tic-tac-toe are *episodic*, with a maximum length to the game (9 for tic-tac-toe and ≈ 6000 for chess).
- The ϵ -greedy algorithm updates episode by episode.
- What about infinite Markov decision processes like robots or long episodes?
 - Rewards received continuously.
 - Not optimal to update episode-by-episode.
- Value function and Q-function learning can update after each step with *Bellman's equations*.

Charu C. Aggarwal
IBM T J Watson Research Center
Yorktown Heights, NY

Value Function Learning and Q-Learning

Neural Networks and Deep Learning, Springer, 2018
Chapter 8.4

Challenges with Long and Infinite Markov Decision Processes

- Previous lecture discusses how value functions can be learned for shorter episodes.
 - Update state-action-value table for each episode with Monte Carlo simulation.
- Effective for games like tic-tac-toe with small episodes.
- What to do with continuous Markov decision processes?

An Infinite Markov Decision Process

- Sequence below is of infinite length (continuous process)

$$s_0 a_0 r_0 s_1 a_1 r_1 \dots s_t a_t r_t \dots$$

- The cumulative reward R_t at time t is given by the discounted sum of the immediate rewards for $\gamma \in (0, 1)$:

$$R_t = r_t + \gamma \cdot r_{t+1} + \gamma^2 \cdot r_{t+2} + \gamma^3 \cdot r_{t+3} \dots = \sum_{i=0}^{\infty} \gamma^i r_{t+i} \quad (1)$$

- Future rewards worth less than immediate rewards ($\gamma < 1$).
- Choosing $\gamma < 1$ is not essential for episodic processes but critical for long MDPs.

Recap of Episodic ϵ -Greedy for Tic-Tac-Toe

- Maintain table of average values of state-action pairs (initialize to small random values).
- Use ϵ -greedy algorithm with table values to simulate moves and create a game.
- **After game:** Update at most 9 entries in the table with new averages, based on the outcomes from $\{\gamma^{r-1}, 0, -\gamma^{r-1}\}$ depending on move number and win/loss.
- Repeat the steps above.

The Bootstrapping Intuition

- Consider a Markov decision process in which we are predicting values (e.g., long-term rewards) at each time-stamp.
 - A partial simulation of the future can improve the prediction at the current time-stamp.
 - This improved prediction can be used as the ground-truth at the current time stamp.
- **Tic-tac-toe:** Parameterized evaluation function for board.
 - After our opponent plays the next move, and board evaluation changes unexpectedly, we go back and correct parameters.
- **Temporal difference learning:** Use difference in prediction caused by partial lookaheads (treated as error for updates).

Example of Chess

- Why is the minimax evaluation of a chess program at 10-ply stronger than that using the 1-ply board evaluation?
 - Because evaluation functions are imperfect (can be strengthened by “cheating” with data from future)!
 - If chess were solved (like checkers today), the evaluation function at any ply would be the same.
 - The minimax evaluation at 10 ply can be used as a “ground truth” for updating a parameterized evaluation function at current position!
- Samuel’s checkers program was the pioneer (called *TD-Leaf* today)
- Variant of idea used by TD-Gammon, *Alpha Zero*.

Q-Learning

- Instead of minimax over a tree, we can use one-step lookahead
- Let $Q(s_t, a_t)$ be a table containing optimal values of state-action pairs (best value of action a_t in state s_t).
- Assume we play tic-tac-toe with ϵ -greedy and $Q(s_t, a_t)$ initialized to random values.
- Instead of Monte Carlo, make following update:
$$Q(s_t, a_t) = r_t + \gamma \max_a Q(s_{t+1}, a) \quad (2)$$
- Update: $Q(s_t, a_t) = Q(s_t, a_t)(1 - \alpha) + \alpha(r_t + \gamma \max_a Q(s_{t+1}, a))$

Why Does this Work?

- Most of the updates we initially make are not meaningful in tic-tac-toe.
 - We started off with random values.
- However, the update of the value of a next-to-terminal state is informative.
- The next time the next-to-terminal state occurs on RHS of Bellman, the update of the next-to-penultimate state will be informative.
- Over time, we will converge to the proper values of all state-action pairs.

SARSA: ϵ -greedy Evaluation

- Let $Q(s_t, a_t)$ be the value of action a_t in state s_t when following the ϵ -greedy policy.
- **An improved estimate** of $Q(s_t, a_t)$ via bootstrapping is $r_t + \gamma Q(s_{t+1}, a_{t+1})$
- Follows from $R_t = \sum_{i=0}^{\infty} \gamma^i r_{t+i} = r_t + \gamma R_{t+1}$
- SARSA: Instead of episodic update, we can update the table containing $Q(s_t, a_t)$ after performing a_t by ϵ -greedy, observing s_{t+1} and then computing a_{t+1} again using ϵ -greedy:


$$Q(s_t, a_t) \leftarrow r_t + \gamma Q(s_{t+1}, a_{t+1}) \quad (3)$$

- Gentler and stable variation: $Q(s_t, a_t) \leftarrow Q(s_t, a_t)(1 - \alpha) + \alpha(r_t + \gamma Q(s_{t+1}, a_{t+1}))$

On-Policy vs Off-Policy Learning

- SARSA: On-policy learning is useful when learning and inference cannot be separated.
 - A robot who continuously learns from the environment.
 - The robot must be cognizant that exploratory actions have a cost (e.g., walking at edge of cliff).
- Q-learning: Off-policy learning is useful when we don't need to perform exploratory component during inference time (have non-zero ϵ during training but set to 0 during inference).
 - Tic-tac-toe can be learned once using Q-learning, and then the model is fixed.

Using Deep Learning




- When the number of states is large, the values $Q(s_t, a_t)$ are predicted from state s_t representation \bar{X}_t rather than tabulated.

$$F(\bar{X}_t, \bar{W}, a) = \hat{Q}(s_t, a) \quad (4)$$

- \bar{X}_t : Previous four screens of pixels in Atari

Specific Details of Convolutional Network



- Same architecture with minor variations was used for all Atari games.

Neural Network Updates for Q-Learning

- The neural network outputs $F(\bar{X}_t, \bar{W}, a_t)$.
- We must wait to observe state \bar{X}_{t+1} and then set up a “ground-truth” value for the output using Bellman’s equations:

$$\text{Bootstrapped Ground-Truth} = r_t + \gamma \max_a F(\bar{X}_{t+1}, \bar{W}, a) \quad (5)$$

- Loss: $L_t = \left\{ \underbrace{[r_t + \gamma \max_a F(\bar{X}_{t+1}, \bar{W}, a)]}_{\text{Treat as constant ground-truth}} - F(\bar{X}_t, \bar{W}, a_t) \right\}^2$
- $\bar{W} \leftarrow \bar{W} + \alpha \left\{ \underbrace{[r_t + \gamma \max_a F(\bar{X}_{t+1}, \bar{W}, a)]}_{\text{Constant ground-truth}} - F(\bar{X}_t, \bar{W}, a_t) \right\} \frac{\partial F(\bar{X}_t, \bar{W}, a_t)}{\partial \bar{W}} \quad (6)$


Neural Network Updates for SARSA

- The neural network outputs $F(\bar{X}_t, \bar{W}, a_t)$.
- We must wait to observe state \bar{X}_{t+1} , simulate a_{t+1} with ϵ -greedy and then set up a “ground-truth” value:

$$\text{Bootstrapped Ground-Truth} = r_t + \gamma F(\bar{X}_{t+1}, \bar{W}, a_{t+1}) \quad (7)$$

- Loss: $L_t = \left\{ \underbrace{[r_t + \gamma F(\bar{X}_{t+1}, \bar{W}, a_{t+1})]}_{\text{Treat as constant ground-truth}} - F(\bar{X}_t, \bar{W}, a_t) \right\}^2$
- $\bar{W} \leftarrow \bar{W} + \alpha \left\{ \underbrace{[r_t + \gamma F(\bar{X}_{t+1}, \bar{W}, a_{t+1})]}_{\text{Constant ground-truth}} - F(\bar{X}_t, \bar{W}, a_t) \right\} \frac{\partial F(\bar{X}_t, \bar{W}, a_t)}{\partial \bar{W}}$ (8)

Value Function Learning



- Instead of outputting values of state-action pairs we can output just values.
- Q-Learning and SARSA can be implemented with this architecture as well.
 - General class of temporal difference learning \Rightarrow Multi-step bootstrapping
 - Can explore a forward-looking tree for arbitrary bootstrapping.

Temporal Difference Learning $TD(0)$

- Value network produces $G(\bar{X}_t, \bar{W})$ and bootstrapped ground truth $= r_t + \gamma G(\bar{X}_{t+1}, \bar{W})$
- Same as SARSA: Observe next state by executing a_t according to current policy
- Loss: $L_t = \left\{ \underbrace{r_t + \gamma G(\bar{X}_{t+1}, \bar{W})}_{\text{"Observed" value}} - G(\bar{X}_t, \bar{W}) \right\}^2$
$$\bar{W} = \bar{W} + \alpha \left\{ \underbrace{[r_t + \gamma G(\bar{X}_{t+1}, \bar{W})]}_{\text{"Observed" value}} - G(\bar{X}_t, \bar{W}) \right\} \frac{\partial G(\bar{X}_t, \bar{W})}{\partial \bar{W}} \quad (9)$$
- Short notation: $\bar{W} \leftarrow \bar{W} + \alpha \delta_t (\nabla G(\bar{X}_t, \bar{W}))$

Bootstrapping over Multiple Steps

- Temporal difference bootstraps only over one time-step.
 - A strategically wrong move will not show up immediately.
 - Can look at n -steps instead of one.
- On-policy looks at single sequence greedily (too weak)
- Off-policy (like Bellman) picks optimal over entire minimax tree (Samuel's checkers program).
- Any optimization heuristic for lookahead-based inference can be exploited.
 - Monte Carlo tree search explores *multiple branches* with upper-bounding strategy \Rightarrow Statistically robust target.

Fixed Window vs Smooth Decay: Temporal Difference Learning $TD(\lambda)$

- Refer to one-step temporal difference learning as $TD(0)$
- Fixed Window n : $\overline{W} \leftarrow \overline{W} + \alpha \delta_t \sum_{k=t-n+1}^t \gamma^{t-k} (\nabla G(\overline{X}_k, \overline{W}))$
- $TD(\lambda)$ corrects past mistakes with discount factor λ when new information is received.

$$\overline{W} \leftarrow \overline{W} + \alpha \delta_t \sum_{k=0}^t (\lambda \gamma)^{t-k} (\nabla G(\overline{X}_k, \overline{W})) \quad (10)$$

- Setting $\lambda = 1$ or $n = \infty$ is equivalent to Monte Carlo methods.
 - Details in book.

Monte Carlo vs Temporal Differences

- Not true that greater lookahead always helps!
 - The value of λ in $TD(\lambda)$ regulates the trade-off between bias and variance.
 - Using small values of λ is particularly advisable if data is limited.
- A temporal difference method places a different value on each position in a single chess game (that depends on the merits of the position).
 - Monte Carlo places a value that depends only on time discounting and final outcome.

Monte Carlo vs Temporal Differences: Chess Example

- Imagine a Monte Carlo rollout of chess game between two agents Alice and Bob.
 - Alice and Bob each made two mistakes but Alice won.
 - Monte Carlo training data does not differentiate between mistakes and good moves.
 - Using temporal differences might see an error after each mistake because an additional ply has *differential* insight about the effect of the move (bootstrapping).
- More data is needed in Monte Carlo rollouts to remove the effect of noise.
- On the other hand, TD(0) might favor learning of end games over openings.

Implications for Other Methods

- Policy gradients often use Monte Carlo rollouts.
 - Deciding the *advantage* of an action is often difficult in games without continuous rewards.
- Value networks are often used in combination with *policy-gradients* in order to design *actor-critic* methods.
 - Temporal differences are used to evaluate the advantage of an action.

Charu C. Aggarwal
IBM T J Watson Research Center
Yorktown Heights, NY


Policy Gradients

Neural Networks and Deep Learning, Springer, 2018
Chapter 8.5

Difference from Value-Based Methods

- Value-based methods like Q-learning attempt to predict the value of an action with a *parameterized value function*.
 - Often coupled with a generic policy (like ϵ -greedy).
- Policy gradient methods estimate the *probability* of each action at each step with the goal of maximizing the overall reward.
- Policy is itself parameterized.

Policy Network vs Q-Network for Atari Game



- Output is *probability* of each action in policy network rather than *value* of each action.

Overview of Approach

- We want to update network to maximize *expected* future rewards ⇒ We need to collect samples of long-term rewards for each simulated action.
 - **Method 1:** Use Monte Carlo policy rollouts to estimate the *simulated* long-term reward after each action.
 - **Method 2:** Use another value network to *model* long-term reward after each action (actor-critic methods).
- Main problem is in setting up a loss function that uses the simulated or modeled rewards to update the parameterized probabilities.

Example: Generating Training Data

- Training chess agent Alice (using pool of human opponents) with reward in $\{+1, 0, -1\}$.
- Consider a Monte Carlo simulation with win for agent Alice.
- Create training points for each board position faced by Alice and each action output a with long-term reward of 1.
 - If discount factor of γ then long-term reward is γ^{r-1} .
- **Backpropagated stochastic gradient ascent:** Somehow need to update neural network from samples of rewards to maximize expected rewards (non-obvious).

Nature of Training Data

- We have board positions together with output action *samples* and long-term rewards of each *sampled* action.
 - We do not have ground-truth *probabilities*.
- So we want to maximize *expected* long-term rewards from *samples* of the probabilistic output.
- How does one compute the gradient of an expectation from samples?

Log Probability Trick

- Let $Q^p(s_t, a)$ be the long-term reward of action a and policy p .
- The log probability trick of REINFORCE relates gradient of expectation to expectation of gradient:

$$\nabla E[Q^p(s_t, a)] = E[Q^p(s_t, a) \nabla \log(p(a))] \quad (11)$$

- $Q^p(s_t, a)$ is estimated by Monte-Carlo roll-out and $\nabla \log(p(a))$ is the log-likelihood gradient from backpropagation in the policy network for sampled action a .

$$\bar{W} \leftarrow \bar{W} + \alpha Q^p(s_t, a) \nabla \log(p(a)) \quad (12)$$

Baseline Adjustments

- Baseline adjustments change the reward into an “advantage” with the use of state-specific adjustments.
 - Subtract some quantity b from each $Q(s_t, a)$
 - Reduces variance of result without affecting bias.
- **State-independent:** One can choose b to be some long-term average of $Q^p(s_t, a)$.
- **State-specific:** One can choose b to be the value $V^p(s_t)$ of state $s_t \Rightarrow$ Advantage is same as temporal difference!

Why Does a State-Specific Baseline Adjustment Make Sense?

- Imagine a self-play chess game in which both sides make mistakes but one side wins.
 - Without baseline adjustments all training points from the game will have long-term rewards that depend only on final results.
 - Temporal difference will capture the *differential* impact of the error made in each action.
 - Gives more refined idea of the specific effect of that move
⇒ Its advantage!


Problems with Monte Carlo Policy Gradients

- Full Monte Carlo simulation is best for episodic processes.
- Actor-critic methods allow online updating by combining ideas in policy gradients and value networks:
 - **Value-based:** The policy (e.g., ϵ -greedy) of the actor is subservient to the critic.
 - **Policy-based:** No notion of critic for value estimation (typical approach is Monte Carlo)
- **Solution:** Create separate neural network for estimating value/advantage.


Actor-Critic Method

- Two separate neural networks:
 - **Actor:** Policy network with parameters $\bar{\Theta}$ that decides actions.
 - **Critic:** Value network or Q-network with parameter \bar{W} that estimates long-term reward/advantage \Rightarrow Advantage is temporal difference.
- The networks are trained simultaneously within an iterative loop.

Actor and Critic



(a) Actor (Decides actions as a probabilistic policy)



(b) Critic (Evaluates advantage in terms of temporal differences)

Steps in Actor-Critic Methods

- Sample the action a_{t+1} at state s_{t+1} using the policy network.
- Use Q-network to compute temporal difference error δ_t at s_t using bootstrapped target derived from value of s_{t+1} .
- **[Update policy network parameters]:** Update policy network using the Q-value of action a_t as its advantage (use temporal difference error for variance reduction).
- **[Update Q-Network parameters]:** Update the Q-network parameters using the squared temporal difference δ_t^2 as the error.
- Repeat the above updates.

Advantages and Disadvantages of Policy Gradients

- Advantages:
 - Work in continuous action spaces.
 - Can be used with stochastic policies.
 - Stable convergence behavior
- Main disadvantage is that they can reach local optima.