

# Gradient Descent Algorithm

Basic PPT with Fundamental  
Equations



# Gradient Descent is often used as black-box tools

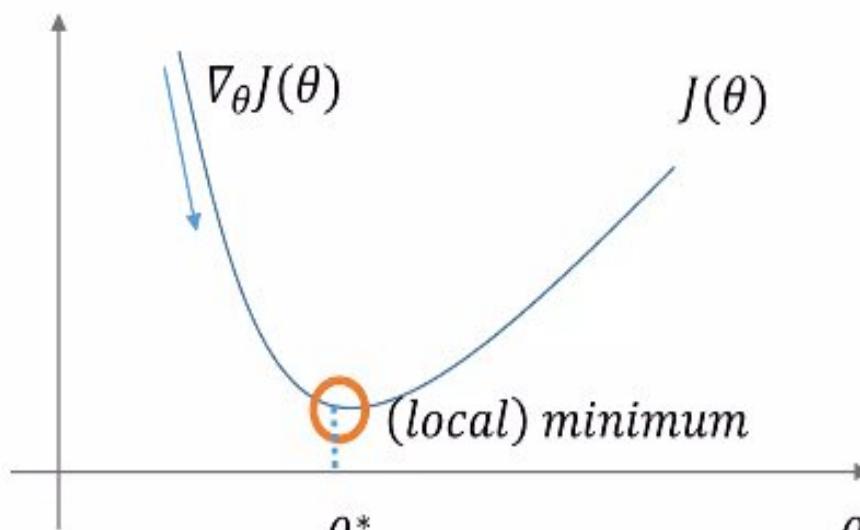
- Gradient descent is popular algorithm to perform optimization of deep learning.
  - Many Deep Learning library contains various gradient descent algorithms.
    - Example : Keras, Chainer, Tensorflow...
- However, **these algorithms often used as black-box tools and many people don't understand their strength and weakness.**
  - We will learn this.

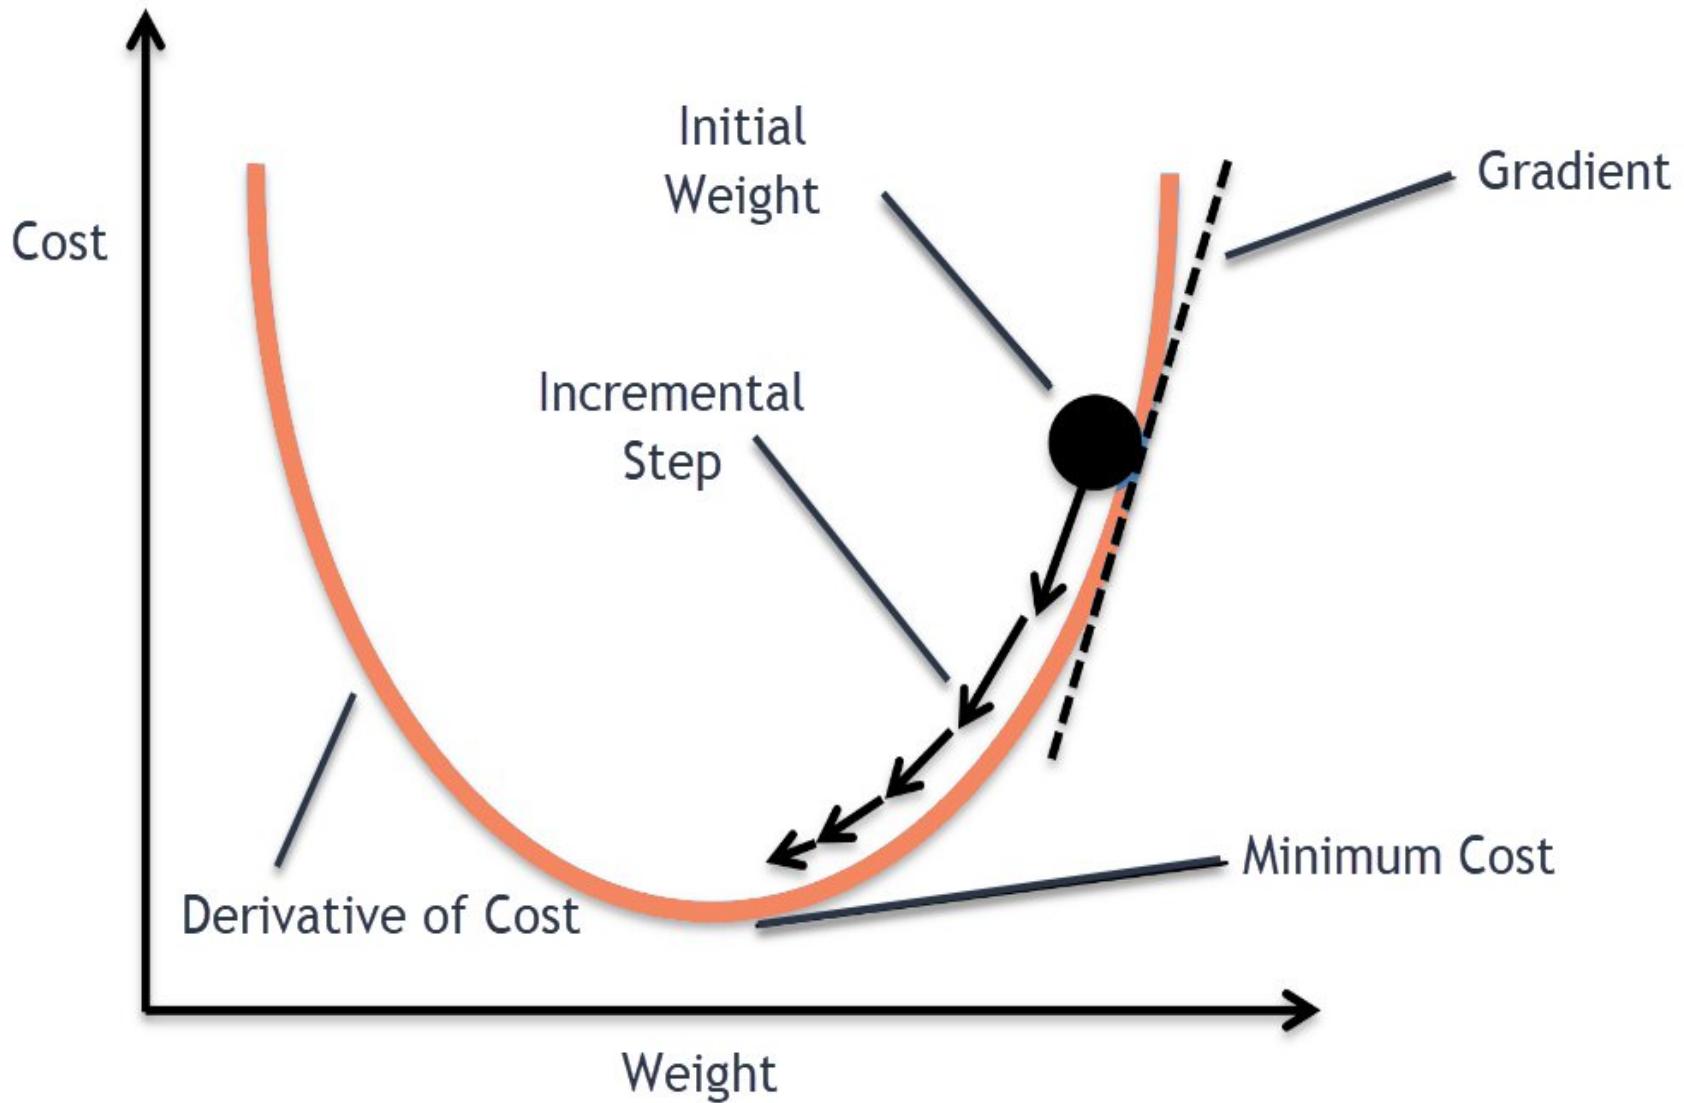
# Gradient Descent

- Gradient descent is a way to minimize an objective function  $J(\theta)$ 
  - $J(\theta)$ : Objective function
  - $\theta \in R^d$ : Model's parameters
  - $\eta$ : Learning rate. This determines the size of the steps we take to reach a (local) minimum.

**Update equation**

$$\theta = \theta - \eta * \nabla_{\theta} J(\theta)$$





# Trade-off

- Depending on the amount of data, they make a trade-off :
  - The **accuracy** of the parameter update
  - The **time** it takes to perform an update.

Method	Accuracy	Time	Memory Usage	Online Learning
Batch gradient descent	○	Slow	High	✗
Stochastic gradient descent	△	High	Low	○
Mini-batch gradient descent	○	Midium	Midium	○

# Batch gradient descent

This method computes the gradient of the cost function  
**with the entire training dataset.**

## Update equation

$$\theta = \theta - \eta * \nabla_{\theta} J(\theta)$$

We need to calculate the gradients for the whole dataset to perform **just one update.**

## Code

```
for i in range(nb_epochs):
    params_grad = evaluate_gradient(loss_function, data, params)
    params = params - learning_rate * params_grad
```

# Batch gradient descent

- Advantage
  - It is guaranteed to converge **to the global minimum for convex error surfaces and to a local minimum for non-convex surfaces.**
- Disadvantages
  - It can be **very slow**.
  - It is intractable for datasets that **do not fit in memory**.
  - It **does not allow** us to update our model **online**.

# Stochastic gradient descent

This method performs a parameter update for **each** training example  $x^{(i)}$  and label  $y^{(i)}$ .

## Update equation

$$\theta = \theta - \eta * \nabla_{\theta} J(\theta; x^{(i)}; y^{(i)})$$

We need to calculate the gradients for the whole dataset to perform **just one update**.

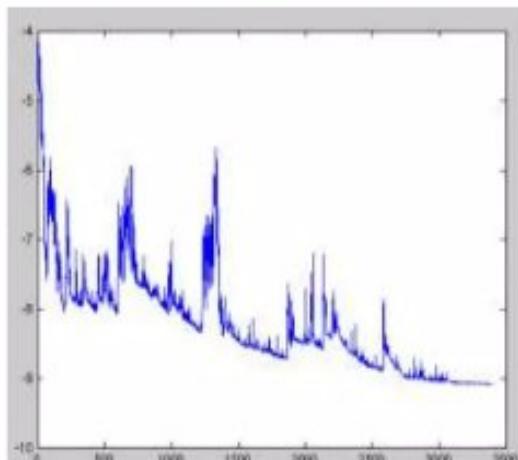
## Code

```
for i in range(nb_epochs):
    np.random.shuffle(data)
    for example in data:
        params_grad = evaluate_gradient(loss_function, example, params)
        params = params - learning_rate * params_grad
```

Note : we shuffle the training data at every epoch

# Stochastic gradient descent

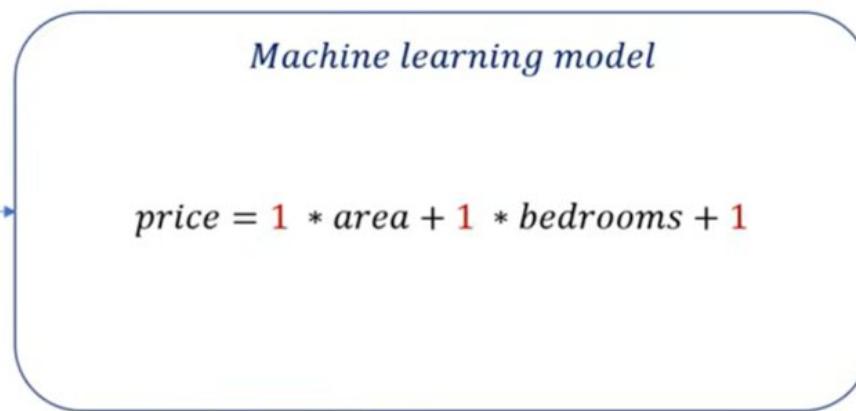
- Advantage
  - It is usually **much faster** than batch gradient descent.
  - It can be **used to learn online**.
- Disadvantages
  - It performs frequent updates with a **high variance** that cause the objective function to fluctuate heavily.



# Example

Initialize  $w_1, w_2$  and  $bias$  to be 1

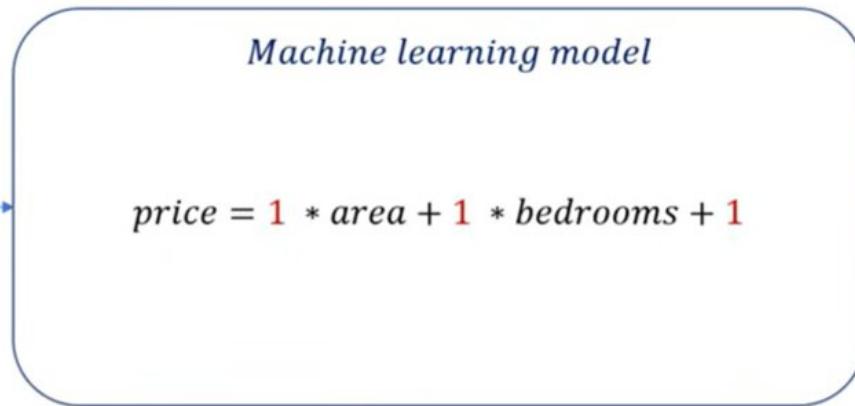
area	bedrooms	price
2600	3	550000
3000	4	565000
3200	3	610000
3600	3	595000
4000	5	760000
4100	6	810000



$$\widehat{price} = 2604$$
$$price = 550000$$

$$error1 = (price - \widehat{price})^2$$

area	bedrooms	price
2600	3	550000
3000	4	565000
3200	3	610000
3600	3	595000
4000	5	760000
4100	6	810000



$$\widehat{price} = 4107$$

$$price = 810000$$



$$error6 = (price - \widehat{price})^2$$

# End of first epoch

*Total Error = error1 + error2 + ... + error6*

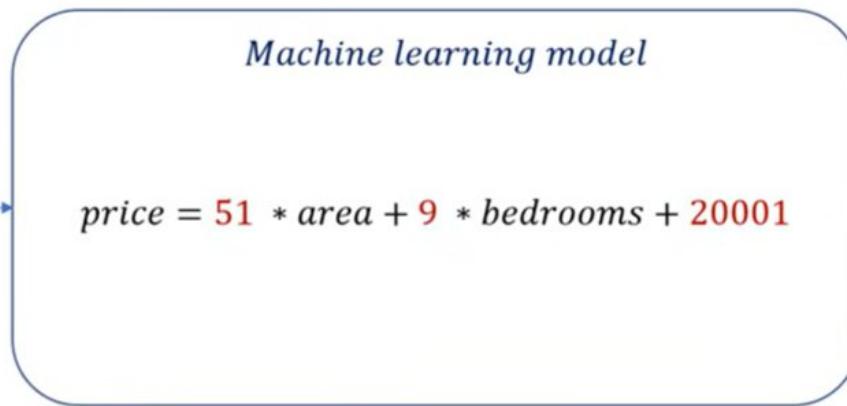
*Mean Squared Error (a.k.a. MSE) =  $\frac{\text{Total Error}}{6}$*

$$w1 = w1 - \text{learning rate} * \frac{\partial(\text{MSE})}{\partial w1}$$
$$w1 = 1 - (-50) = 51$$

$$w2 = w2 - \text{learning rate} * \frac{\partial(\text{MSE})}{\partial w2}$$
$$w2 = 1 - (-8) = 9$$

$$b = b - \text{learning rate} * \frac{\partial(\text{MSE})}{\partial b}$$
$$\text{bias} = 1 - (-20000) = 20001$$

area	bedrooms	price
2600	3	550000
3000	4	565000
3200	3	610000
3600	3	595000
4000	5	760000
4100	6	810000



$$\widehat{price} = 229155$$
$$price = 550000$$

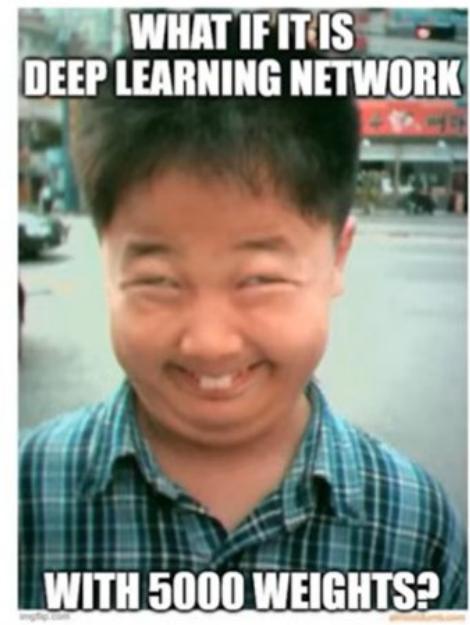
$$error6 = (price - \widehat{price})^2$$

**End of second epoch**

area	bedrooms	price
2600	3	550000
3000	4	565000
3200	3	610000
3600	3	595000
...	...	...
4100	6	810000

*10 million samples*

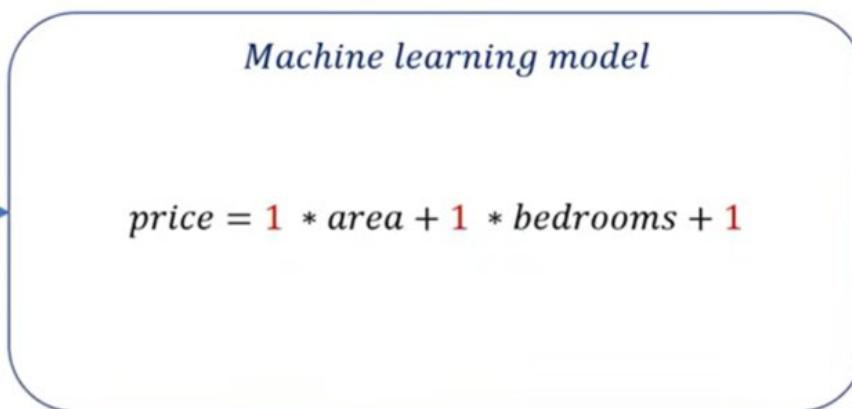
- 1) To find cumulative error for first round (epoch) now we need to do a forward pass for **10 million samples**
- 2) We have **2 features** (area and bedroom). This requires finding **20 million derivatives**



## 1. Randomly pick single data training sample

area	bedrooms	price
2600	3	550000
3000	4	565000
3200	3	610000
3600	3	595000
...	...	...
4100	6	810000

10 million samples



$$error = (price - \widehat{price})^2$$

## 2. Adjust weights

$$w_1 = w_1 - \text{learning rate} * \frac{\partial(\text{error})}{\partial w_1} \quad w_1 = 1 - (-13) = 14$$

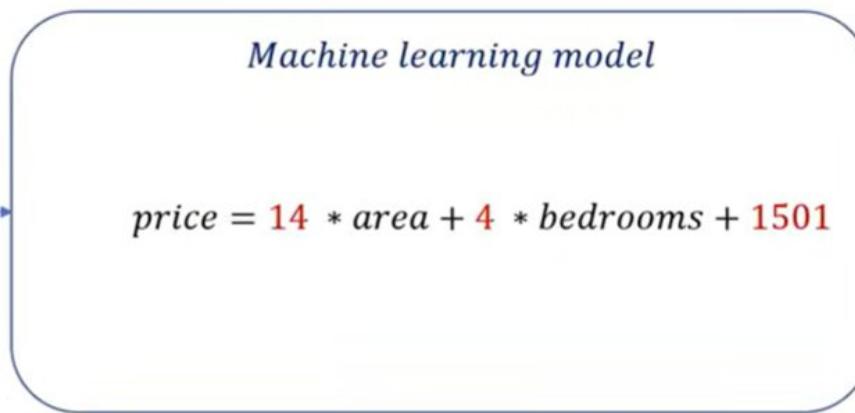
$$w_2 = w_2 - \text{learning rate} * \frac{\partial(\text{error})}{\partial w_2} \quad w_2 = 1 - (-3) = 4$$

$$b = b - \text{learning rate} * \frac{\partial(\text{error})}{\partial b} \quad \text{bias} = 1 - (-1500) = 1501$$

### 3. Again randomly pick a training sample

area	bedrooms	price
2600	3	550000
3000	4	565000
3200	3	610000
3600	3	595000
...	...	...
4100	6	810000

10 million samples



$$\widehat{price} = 3204$$

$$price = 37913$$



$$error = (price - \widehat{price})^2$$

#### *4. Again adjust weights*

$$w_1 = w_1 - \text{learning rate} * \frac{\partial(\text{error})}{\partial w_1} \quad w_1 = 14 - (-100) = 114$$

$$w_2 = w_2 - \text{learning rate} * \frac{\partial(\text{error})}{\partial w_2} \quad w_2 = 4 - (-12) = 16$$

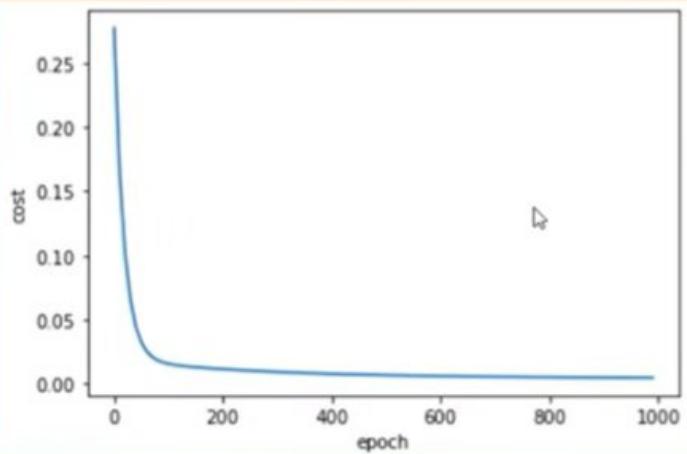
$$b = b - \text{learning rate} * \frac{\partial(\text{error})}{\partial b} \quad \text{bias} = 1501 - (-2001) = 3502$$



### Batch Gradient Descent

Use **all** training samples for one forward pass and then adjust weights

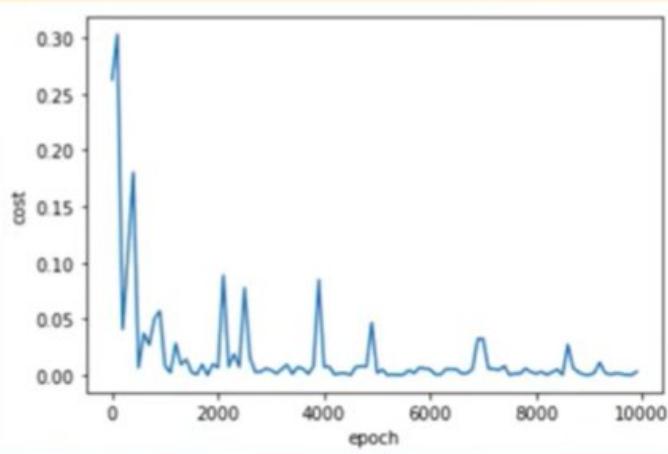
Good for small training set

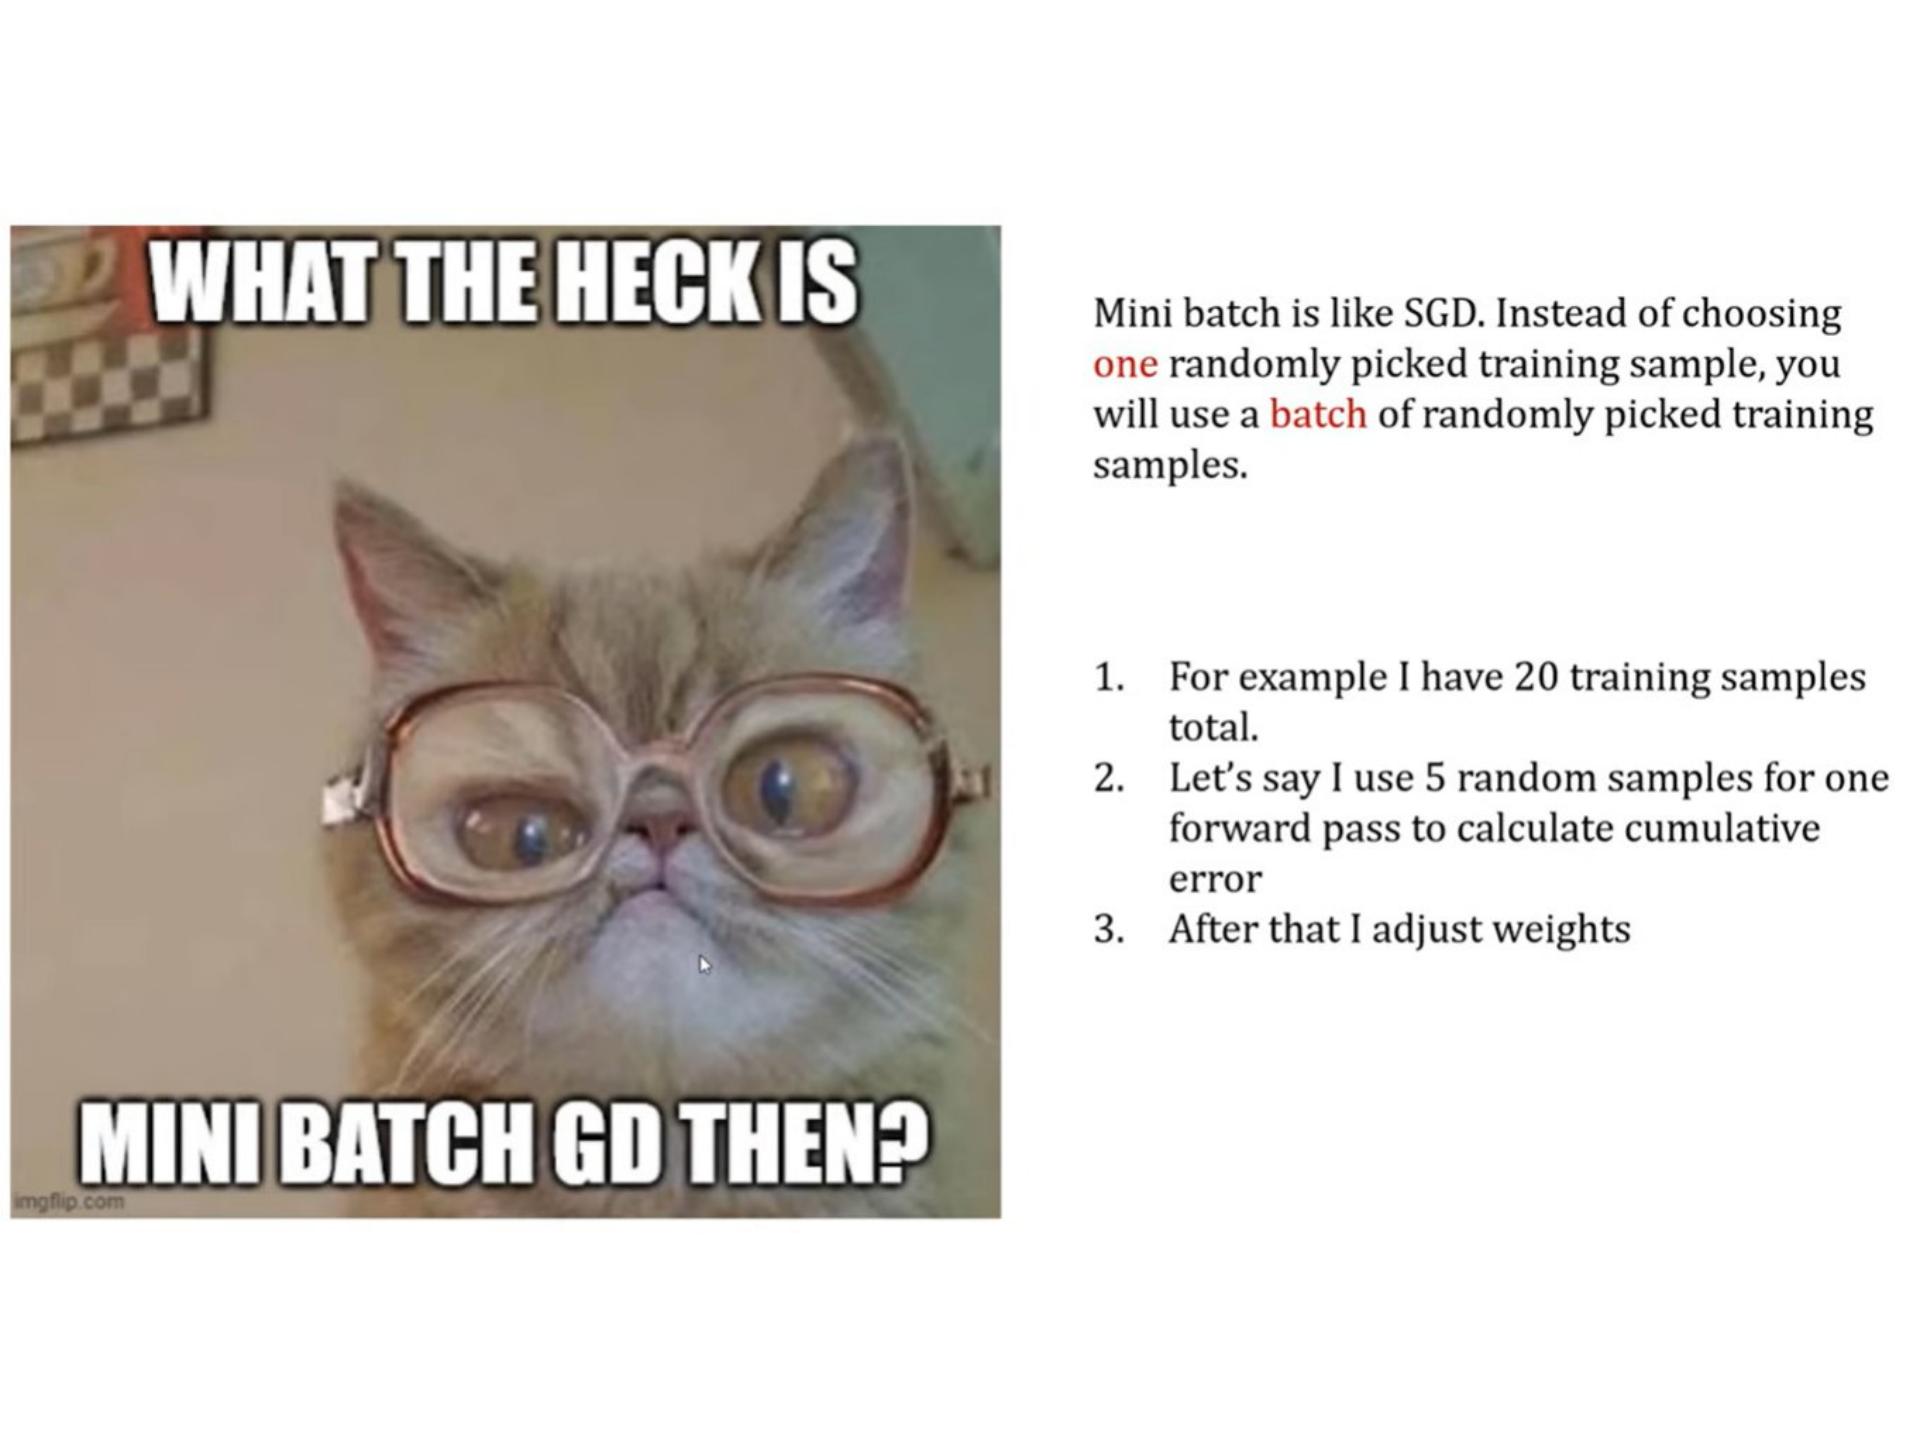


### Stochastic Gradient Descent (SGD)

Use **one** (randomly picked) sample for a forward pass and then adjust weights

Good when training set is very big and we don't want too much computation





**WHAT THE HECK IS**

**MINI BATCH GD THEN?**

Mini batch is like SGD. Instead of choosing **one** randomly picked training sample, you will use a **batch** of randomly picked training samples.

1. For example I have 20 training samples total.
2. Let's say I use 5 random samples for one forward pass to calculate cumulative error
3. After that I adjust weights

# Gradient Descent Comparison Table (Updated)

Feature	Batch Gradient Descent (BGD)	Stochastic Gradient Descent (SGD)	Mini-Batch Gradient Descent (MBGD)
Gradient Computed On	Entire dataset	One sample at a time	A small subset (batch) of data (e.g., 32, 64, 128)
Updates per Epoch	1 update	N updates (N = no. of samples)	N / batch_size updates
Speed	Slow	Fast	Faster than BGD, slower than SGD
Memory Requirement	High	Very low	Moderate
Convergence	Smooth & stable	Noisy, oscillates	Balanced, smooth but faster
Accuracy	High (best gradient estimate)	Lower due to noise	High (close to BGD)
Suitable For	Small–medium datasets	Very large datasets, online learning	Most practical ML problems
Pros	Stable convergence, full-data accuracy	Very fast updates, low memory	Good trade-off between speed & accuracy
Cons	Computationally expensive	Noisy updates, may overshoot	Requires tuning batch size